

2 Project Objectives

The objective of this project is to build TCP from scratch. The project is divided into two tasks: reliable data transfer and congestion control. Both tasks have individual due dates and are interdependent.

Since this is a group project, you are required to find a (1) partner to work with. Once you have settled on a partner, you need to update the [Excel sheet](#) shared with you in the lab session before the deadline specified in the submission details and policy section. Note that if the deadline for forming a group is missed, a 10% penalty will be imposed.

Before beginning the project, make sure to read the **entire** project description carefully. This will help you understand the provided functionality and the features you're expected to implement.

3 Task 1: Simplified TCP Sender/Receiver

For the first task, the goal is to create a "Reliable Data Transfer" protocol that aligns with the description in section 3.5.4 of the textbook. The objective is to develop a simplified TCP sender and receiver that can manage packet losses and retransmissions.

The implementation of the protocol requires the following functionalities:

- Sending packets to the network based on a fixed sending window size, as illustrated in Figure 3.33 in the textbook. The window size is set to 10 packets.
- Sending cumulative acknowledgments from the receiver and determining how to respond to them at the sender.
- Using a single retransmission timer to handle packet loss and retransmission.

In short, the approach used for implementing the first task is similar to the Go-Back-N protocol. Both use cumulative acknowledgments and a single timeout timer for the oldest unacknowledged packet. However, in this approach, only the packet with the smallest sequence number in the window is retransmitted upon a timeout, as opposed to the entire window in the Go-Back-N protocol.

At the receiver, out-of-order packets must be buffered and cumulative acknowledgments are sent for out-of-order packets. When a packet is lost, the packet must be retransmitted by the sender using a retransmission timeout timer (RTO) with a fixed timeout value that is suitable for the emulated network scenario using MahiMahi. When an ACK is received that acknowledges a transmitted packet, the retransmission timer is restarted so that it will expire after the fixed timeout value.

The sender should terminate upon successfully transmitting the entire file, including receiving an acknowledgment for the very last packet.

We have provided you with a simple (stop-and-wait) starter-code that consists of the following:

- rdt receiver: this holds the implementation of a simple reliable data transfer protocol (rdt) receiver, similar to Figure 3.14 in the textbook.

- rdt sender: this holds the implementation of a simple reliable data transfer protocol (rdt) sender, similar to Figure 3.15 in the textbook.
- Channel traces for emulating different network conditions

The simple rdt protocol is implemented on top of the UDP transport protocol. During the lab session, the TA showed you how to use the network emulator MahiMahi to test your sender and receiver functionality in an emulated network environment.

3.1 Sliding Window

The sender and the receiver have to maintain a sliding window, as shown in Figure 1.




Figure 1: Sliding Window

The sender slides the window forward when it receives an ACK for a packet with the lowest sequence number in the sliding window. There is a sequence number associated with each packet and the following constraints are valid for the sender:

1. $\text{LastPacketAked} \leq \text{LastPacketSent}$
2. $\text{LastPacketSent} \leq \text{LastPacketAvailable}$
3. $\text{LastPacketSent} - \text{LastPacketAked} \leq \text{WindowSize}$
4. Packets between LastPacketAked and $\text{LastPacketAvailable}$ must be “buffered”. You can either implement this by buffering the packets or by being able to regenerate them from the data file.

After the sender sends a data packet, it starts a timer (if it is not already running) and waits for a certain period to receive an acknowledgment for the packet. Meanwhile, whenever the receiver receives a packet, it responds with an ACK for the `NextPacketExpected` (cumulative ACK). If the receiver receives an out-of-order packet, it should buffer the packet and send a duplicate ACK. For example, upon receiving a packet with sequence number = 100, the reply would be “ACK 101”, but only if all packets with sequence numbers less than 100 have already been received. These ACKs are called cumulative ACKs. The sender has two ways to know if the packets it sent did not reach the receiver: either a timeout occurred, or the sender received “duplicate ACKs”.

If the sender sent a packet and did not receive an ACK for it before the timer expired, it retransmits the packet. If the sender sent a packet and received duplicate ACKs, it knows that the next expected packet

(at least) was lost. To avoid confusion from reordering, a sender counts a packet lost only after 3 duplicate ACKs in a row.

4 Task 2: TCP Congestion Control

Task 2 involves implementing a congestion control protocol for the sender and receiver developed in Task 1, similar to TCP Tahoe. The congestion control protocol should include the following features:

- Slow-start
- Congestion avoidance
- Fast retransmit (without fast recovery)

The next subsections detail the requirements of the assignment. This high-level outline roughly mirrors the order in which you should implement the functionality. For further details, please refer to the lecture slides or the textbook.

4.1 Retransmission Timer

Task 1 assumed a fixed value for the retransmission timeout timer (RTO). In task 2, the duration of the timeout timer should be determined based on the Round Trip Time (RTT) estimator explained in the lecture and in RFC 2988. The RTO should include Karn's algorithm and also implement an "exponential backoff" for successive timeouts for the same segment. The following values should be used to initialize the RTT estimation parameters:

- RTT: 0 seconds
- RTO: 3 seconds

The upper bound for the RTO is 240 seconds and may be used to provide an upper bound for the exponential backoff.

4.2 Congestion Control

Broadly speaking, the idea of TCP's congestion control is to determine how much capacity is available in the network, so it knows how many packets it can safely have "in-flight" at the same time. Once the sender has this many packets in transit, it uses the arrival of an ACK as a signal that one of its packets has left the network, and it is therefore safe to insert a new packet into the network without adding to the level of congestion.

TCP is considered to be "self-clocking" because it uses acknowledgments to regulate the rate at which packets are transmitted. In addition, TCP incorporates a congestion control mechanism that comprises several algorithms, including Slow Start, Congestion Avoidance, Fast Retransmit, and Fast Recovery. You can read more about these mechanisms in the textbook Section 3.7. In the first part of the project, the window size was fixed to 10 packets. The task of the second part is to dynamically determine the ideal congestion window size (CWND) of the sender.

Slow-start: When a new connection is established with a host, the CWND is initialized to one packet. Each time an ACK is received, CWND is increased by one packet. The sender keeps increasing CWND until the first packet loss is detected or until CWND reaches the value ssthresh (Slow-start threshold), after which it enters Congestion Avoidance (see below). For a new connection, the ssthresh is set to a very large value, i.e. 64 packets. If a packet is lost in the slow-start phase, the sender sets ssthresh to $\max(\text{CWND}/2, 2)$ in case the client later returns to Slow-start again.

Congestion Avoidance: After the slow-start phase, the CWND is slowly increased until a packet loss occurs. The increase of the CWND should be at most one packet per round-trip time (regardless how many ACKs are received in that RTT). That is, when the sender receives an ACK, it usually increases the CWND by a fraction equal to $1/\text{CWND}$. You may notice here that you need to use a float variable for the CWND, however when you send data you are always going to take the floor of the CWND. As soon as the entire window is acknowledged, only then these fractions would sum to a 1.0 and as a result the CWND would then have increased by 1 packet. This is in contrast to Slow-start where CWND is incremented for each ACK. Recall that when the sender receives 3 duplicate ACKs, it can be assumed that the packet with sequence number == acknowledgment number was lost, even if a timeout has not occurred yet. This process is called Fast Retransmit.

Fast Retransmit: Once a Fast Retransmit occurs, the ssthresh is set to $\max(\text{CWND}/2, 2)$. CWND is then set to 1 and the Slow-start process starts again. Please note that implementing the fast recovery mechanism is not required for this project.

4.3 Graphing CWND

As part of the sender implementation, you are required to generate an output file called "CWND.csv", which tracks how the congestion window (CWND) varies over time. This file serves both as a debugging tool for testing your code and as a means for us to grade your submission. The file should include the time when the window changed and the current value of the CWND. To facilitate grading, you should also update the plotting script (plot.py) to graphically depict the evolution of the CWND over time.

5 Further Implementation Details

Please note the following implementation details for **both** tasks:

- The handshake procedures of TCP (connection establishment and termination), flow control, and fast recovery do not need to be implemented.
- Please note that the setup of sender and receiver in this project does not reflect a real world scenario. For simplicity, in this project, the receiver must be started first and the sender second. Once the sender is started, it will begin sending the file to the receiver. In a real world scenario, the sender would act as the server and the receiver as the client. Furthermore, the client would request a file from the server (typically using the application layer) and wait for the incoming file.
- To avoid overwriting the sending file with the receiving file, ensure that you either give the output file at the receiver a different name or run the sender and receiver code in two different folders.

- The sliding window should have a maximum size of a 32 bit integer and wraps around, as discussed in section 3.4.3 in the textbook.
- "Delayed ACKs" at the receiver do not need to be implemented.
- Sequence and acknowledgment numbers are based on bytes, rather than on discrete packet numbers.
- Your submission **must** be working successfully on a Linux Operating System (either the GUI of the remote Linux Server shown in the lab or using VirtualBox with an Ubuntu image).

6 Grading

Description	Score (/20)
Task 1: Extending the sender to send 10 packets	2
Task 1: Buffering of out-of order packets at the receiver	2
Task 1: Properly sending and handling ACKs	1
Task 1: Retransmissions of lost packets	1.5
Task 1: Properly receiving the exact file on the receiver (no errors)	1.5
Task 2: Slow-start implementation	2
Task 2: Congestion avoidance implementation	3
Task 2: Fast retransmit implementation	2
Task 2: Retransmission timeout timer	2
Task 2: Correct CWND recording and plotting	2
Coding style and usage of meaningful comments	1

7 Submission Details and Policy

Submission Deadlines:

1. Group formation due date: February 17, 2025 (-10% penalty if not met)
2. Task 1 due date: February 28, 2025 (40% of grade)
3. Task 2 due date: March 26, 2025 (60% of grade)

Submission Format and System: You can directly submit **all** your files as a zip file on Brightspace (<https://brightspace.nyu.edu/>). Due to technical limitations, submissions via email are not accepted.

Late Submissions: Late submissions will be penalized by 10% per 24 hours, with a maximum of 3 days late. In case of a late submission, please upload your zip file to Brightspace.